

Lecture 12 – Multiple Linear Regression and Feature Engineering



DSC 40A, Winter 2024

Announcements

- ▶ No homework due this week.
 - ▶ Homework 4 will be out this Friday (Feb. 9) after Midterm 1.
- ▶ Homework 4 due date extend to next Friday (Feb. 16)

Midterm 1 is Friday during lecture

- ▶ [Formula sheet](#) will be provided for you. No other notes.
- ▶ No calculators. This implies no crazy calculations.
- ▶ Assigned seats will be posted on Course Website and Campuswire.
- ▶ We will not answer questions during the exam. State your assumptions if anything is unclear.
- ▶ The exam will include long-answer homework-style questions, as well as short-answer questions such as True/False or filling in a numerical answer.
- ▶ The exam covers Homeworks 1 through 3 plus additional contents in linear algebra and multiple linear regression.

Midterm study strategy

- ▶ Review the written solutions to previous homeworks and groupworks.
- ▶ Identify which concepts are still uncertain. Re-watch podcasts, post on Campuswire, come to office hours, use resources [on course website](#), watch Janine's lecture videos.
- ▶ Work through past exams [on course website](#) and the posted mock exam.
- ▶ Study in groups.
- ▶ Summarize key facts and formulas.
- ▶ Some problems are easier, some problems are harder.

Some Tips About Midterm

- ▶ Understand the derivations we did in lecture.
 - ▶ There will be derivation problems in the midterm, but no long derivation.
 - ▶ Understand the derivation I did in lecture and some methods I used
- ▶ Be able to perform simple algebra, calculus and linear algebra computation
 - ▶ Example: calculating matrix multiplication.
- ▶ Read each question carefully.
 - ▶ Example: using formal definition to prove convexity vs. using any method to prove convexity.

Extra Credit Opportunity

- ▶ The last problem on HW4 will be a class-wide competition on finding energies for High-Purity Germanium Detector waveforms
 - ▶ I'll explain what this is on next Monday's lecture, after the exam.
- ▶ Top predictions will get extra credit on Midterm 1.
- ▶ More detail next Monday

Agenda

- ▶ Incorporating multiple features.
- ▶ Interpreting parameters.
- ▶ Feature engineering.

Incorporating multiple features

Last time

- ▶ We minimized the mean squared error for the prediction rule $H(x) = w_0 + w_1x$, which was

$$R_{sq}(\vec{w}) = \frac{1}{n} ||\vec{y} - X\vec{w}||^2$$

- ▶ We found that the minimizing \vec{w} satisfies the **normal equations**, $X^T X \vec{w} = X^T \vec{y}$.
 - ▶ If $X^T X$ is invertible, the solution is:

$$\vec{w}^* = (X^T X)^{-1} X^T \vec{y}$$

- ▶ These same normal equations can be used to solve the **multiple linear regression** problem, where we use multiple features to predict an outcome. We simply need to adjust the design matrix X .

Multiple linear regression example

- ▶ We're want to fit a **linear** prediction rule with two features:

$$H(\text{experience, GPA}) = w_0 + w_1(\text{experience}) + w_2(\text{GPA})$$

- ▶ Collect data for each of n people:

Person #	Experience	GPA	Salary
1	3	3.7	85,000
2	6	3.3	95,000
3	10	3.1	105,000

- ▶ We represent each person with a **feature vector**:

$$\vec{x}_1 = \begin{bmatrix} 3 \\ 3.7 \end{bmatrix}, \quad \vec{x}_2 = \begin{bmatrix} 6 \\ 3.3 \end{bmatrix}, \quad \vec{x}_3 = \begin{bmatrix} 10 \\ 3.1 \end{bmatrix}$$

Prediction rule form determines design matrix

- When our prediction rule is

$$H(\text{experience}, \text{GPA}) = w_0 + w_1(\text{experience}) + w_2(\text{GPA}),$$

the hypothesis vector $\vec{h} \in \mathbb{R}^n$ can be written

$$\vec{h} = \begin{bmatrix} H(\text{experience}_1, \text{GPA}_1) \\ H(\text{experience}_2, \text{GPA}_2) \\ \dots \\ H(\text{experience}_n, \text{GPA}_n) \end{bmatrix} = \begin{bmatrix} 1 & \text{experience}_1 & \text{GPA}_1 \\ 1 & \text{experience}_2 & \text{GPA}_2 \\ \dots & \dots & \dots \\ 1 & \text{experience}_n & \text{GPA}_n \end{bmatrix} \begin{bmatrix} w_0 \\ w_1 \\ w_2 \end{bmatrix}$$

- Notice that the rows of the design matrix are the (transposed) feature vectors, with an additional 1 in front.

Notation for multiple linear regression

- ▶ We will need to keep track of multiple¹ features for every individual in our data set.
- ▶ As before, subscripts distinguish between individuals in our data set. We have n individuals (or **training examples**).
- ▶ Superscripts distinguish between features.² We have d features.
 - ▶ experience = $x^{(1)}$
 - ▶ GPA = $x^{(2)}$

¹In practice, we might use hundreds or even thousands of features.

²Think of them as new variable names, such as new letters.

Augmented feature vectors

- The **augmented feature vector** $\text{Aug}(\vec{x})$ is the vector obtained by adding a 1 to the front of feature vector \vec{x} :

$$\vec{x} = \begin{bmatrix} x^{(1)} \\ x^{(2)} \\ \vdots \\ x^{(d)} \end{bmatrix} \quad \text{Aug}(\vec{x}) = \begin{bmatrix} 1 \\ x^{(1)} \\ x^{(2)} \\ \vdots \\ x^{(d)} \end{bmatrix} \quad \vec{w} = \begin{bmatrix} w_0 \\ w_1 \\ w_2 \\ \vdots \\ w_d \end{bmatrix}$$

- Then, our prediction rule is

$$\begin{aligned} H(\vec{x}) &= w_0 + w_1 x^{(1)} + w_2 x^{(2)} + \dots + w_d x^{(d)} \\ &= \vec{w} \cdot \text{Aug}(\vec{x}) \end{aligned}$$

The general problem

- ▶ We have n data points (or **training examples**):
 $(\vec{x}_1, y_1), \dots, (\vec{x}_n, y_n)$ where each \vec{x}_i is a feature vector of d features:

$$\vec{x}_i = \begin{bmatrix} x_i^{(1)} \\ x_i^{(2)} \\ \dots \\ x_i^{(d)} \end{bmatrix}$$

- ▶ We want to find a good linear prediction rule:

$$\begin{aligned} H(\vec{x}) &= w_0 + w_1 x^{(1)} + w_2 x^{(2)} + \dots + w_d x^{(d)} \\ &= \vec{w} \cdot \text{Aug}(\vec{x}) \end{aligned}$$

The general solution

- Use design matrix

$$X = \begin{bmatrix} 1 & x_1^{(1)} & x_1^{(2)} & \dots & x_1^{(d)} \\ 1 & x_2^{(1)} & x_2^{(2)} & \dots & x_2^{(d)} \\ \dots & \dots & \dots & \dots & \dots \\ 1 & x_n^{(1)} & x_n^{(2)} & \dots & x_n^{(d)} \end{bmatrix} = \begin{bmatrix} \text{Aug}(\vec{x}_1)^T \\ \text{Aug}(\vec{x}_2)^T \\ \dots \\ \text{Aug}(\vec{x}_n)^T \end{bmatrix}$$

and observation vector to solve the **normal equations**

$$X^T X \vec{w}^* = X^T \vec{y}$$

to find the optimal parameter vector.

Terminology for parameters

- ▶ With d features, \vec{w} has $d + 1$ entries.
- ▶ w_0 is the **bias**, also known as the **intercept**.
- ▶ w_1, \dots, w_d each give the **weight**, i.e. **coefficient**, of a feature.

$$H(\vec{x}) = w_0 + w_1 x^{(1)} + \dots + w_d x^{(d)}$$

Interpreting parameters

Example: predicting sales

- ▶ For each of 26 stores, we have:
 - ▶ net sales,
 - ▶ square feet,
 - ▶ inventory,
 - ▶ advertising expenditure,
 - ▶ district size, and
 - ▶ number of competing stores.
- ▶ Goal: predict net sales given these features
- ▶ To begin:

$$H(\text{square feet, competitors}) = w_0 + w_1(\text{square feet}) + w_2(\text{competitors})$$

Example: predicting sales

$$H(\text{square feet, competitors}) = w_0 + w_1(\text{square feet}) + w_2(\text{competitors})$$

Discussion Question

What will be the sign of w_1^* and w_2^* ?

- a) $w_1^* = +$, $w_2^* = -$
- b) $w_1^* = +$, $w_2^* = +$
- c) $w_1^* = -$, $w_2^* = -$
- d) $w_1^* = -$, $w_2^* = +$

Example: predicting sales

$$H(\text{square feet, competitors}) = w_0 + w_1(\text{square feet}) + w_2(\text{competitors})$$

Discussion Question

What will be the sign of w_1^* and w_2^* ?

- a) $w_1^* = +$, $w_2^* = -$
- b) $w_1^* = +$, $w_2^* = +$
- c) $w_1^* = -$, $w_2^* = -$
- d) $w_1^* = -$, $w_2^* = +$

Let's try it out ourselves. [Follow along here.](#)

Which features are most “important”?

Discussion Question

Which feature has the greatest effect on the outcome?

- a) square feet: $w_1^* = 16.202$
- b) competitors: $w_2^* = -5.311$
- c) inventory: $w_2^* = 0.175$
- d) advertising: $w_3^* = 11.526$
- e) district size: $w_4^* = 13.580$

Which features are most “important”?

- ▶ The most important feature is **not necessarily** the feature with largest weight.
- ▶ Features are measured in different units, scales.
 - ▶ Suppose I fit one prediction rule, H_1 , with sales in dollars, and another prediction rule, H_2 , with sales in thousands of dollars.
 - ▶ Sales is just as important in both prediction rules.
 - ▶ But the weight of sales in H_1 will be 1000 times smaller than the weight of sales in H_2 .
 - ▶ Intuitive explanation: $5 \times 45000 = (5 \times 1000) \times 45$.
- ▶ **Solution:** before doing regression, **standardize** each feature, i.e. convert each feature to standard units.

Standard units

- ▶ Recall: to convert a feature x_1, x_2, \dots, x_n to standard units, we use the formula

$$x_i \text{ in standard units} = \frac{x_i - \bar{x}}{\sigma_x}$$

- ▶ Example: 1, 7, 7, 9
 - ▶ Mean:
 - ▶ Standard deviation:
 - ▶ Standardized data:


Standard units for multiple linear regression

- ▶ The result of standardizing each feature (separately!) is that the units of each feature are on the same scale.
 - ▶ There's no need to standardize the outcome (net sales), since it's not being compared to anything.
- ▶ Then, solve the normal equations. The resulting $w_0^*, w_1^*, \dots, w_d^*$ are called the **standardized regression coefficients**.
- ▶ Standardized regression coefficients can be directly compared to one another.

Let's try it out in our demo notebook.

Feature engineering

MPG vs. Horsepower



Question: Would a linear prediction rule work well on this dataset?

A quadratic prediction rule

- ▶ It looks like there's some sort of quadratic relationship between horsepower and mpg in the last scatter plot. We want to try and fit a prediction rule of the form

$$H(x) = w_0 + w_1 x + w_2 x^2$$

- ▶ Note that while this is quadratic in horsepower, it is **linear in the parameters!**
- ▶ We can do that, by choosing our two “features” to be x_i and x_i^2 , respectively.
 - ▶ In other words, $x_i^{(1)} = x_i$ and $x_i^{(2)} = x_i^2$.
 - ▶ More generally, we can create new features out of existing features.

A quadratic prediction rule

- ▶ Desired prediction rule: $H(x) = w_0 + w_1x + w_2x^2$.
- ▶ The resulting design matrix looks like this:

$$X = \begin{bmatrix} 1 & x_1 & x_1^2 \\ 1 & x_2 & x_2^2 \\ \dots & & \\ 1 & x_n & x_n^2 \end{bmatrix}$$

- ▶ To find optimal parameter vector \vec{w}^* : solve the **normal equations**!

$$X^T X w^* = X^T y$$

More examples

- ▶ What if we want to use a prediction rule of the form
$$H(x) = w_0 + w_1 x + w_2 x^2 + w_3 x^3?$$
- ▶ What if we want to use a prediction rule of the form
$$H(x) = w_1 \frac{1}{x^2} + w_2 \sin x + w_3 e^x?$$

Feature engineering

- ▶ More generally, we can create new features out of existing information in our dataset. This process is called **feature engineering**.
 - ▶ In this class, feature engineering will mostly be restricted to creating non-linear functions of existing features (as in the previous example).
 - ▶ In the future you'll learn how to do other things, like encode categorical information.

Summary

Summary

- ▶ The normal equations can be used to solve the **multiple linear regression** problem, where we use multiple features to predict an outcome.
- ▶ We can interpret the parameters as weights. The signs of weights give meaningful information, but we can only compare weights if our features are standardized.
- ▶ We can create non-linear features out of existing features. This process is called **feature engineering**.
 - ▶ A prediction rule only needs to be a **linear function of the parameters** for us to use linear regression. It does not need to be a linear function of the features.