

## Lecture 13 – Feature Engineering, Clustering



DSC 40A, Winter 2024

## Midterm 1 is Friday during lecture

- ▶ [Formula sheet](#) will be provided for you. No other notes.
- ▶ No calculators. This implies no crazy calculations.
- ▶ Assigned seats will be posted on Course Website and Campuswire.
- ▶ We will not answer questions during the exam. State your assumptions if anything is unclear.
- ▶ The exam will include long-answer homework-style questions, as well as short-answer questions such as True/False or filling in a numerical answer.
- ▶ The exam covers Lecture 1 to Lecture 12 (HW1-3 + Linear Algebra and Multiple Linear Regression).

## Midterm study strategy

- ▶ Look at annotated lecture notes.
- ▶ Review the written solutions to previous homeworks and groupworks.
- ▶ Identify which concepts are still uncertain. Re-watch podcasts, post on Campuswire, come to office hours, use resources [on course website](#), watch Janine's lecture videos.
- ▶ Work through past exams [on course website](#) and the posted mock exam.
- ▶ Study in groups.
- ▶ Summarize key facts and formulas.

## Some Tips About Midterm

- ▶ Understand the derivations we did in lecture.
  - ▶ There will be derivation problems in the midterm, but no long derivation.
  - ▶ Understand the derivation I did in lecture and some methods I used
- ▶ Be able to perform simple algebra, calculus and linear algebra computation
  - ▶ Example: calculating matrix multiplication.
- ▶ Read each question carefully.
  - ▶ Example: using formal definition to prove convexity vs. using any method to prove convexity.
- ▶ Some problems are easier, some problems are harder.

## Extra Credit Opportunity

- ▶ The last problem on HW4 will be a class-wide competition on finding energies for High-Purity Germanium Detector waveforms
  - ▶ I'll explain what this is on next Monday's lecture, after the exam.
- ▶ Top predictions will get extra credit on Midterm 1.
- ▶ More detail next Monday


# Agenda

- ▶ Feature engineering.
- ▶ Taxonomy of machine learning.
- ▶ Clustering.

## Feature engineering

# Last time: Cars

MPG vs. Horsepower



**Question:** Would a linear prediction rule work well on this dataset?



## A quadratic prediction rule

- ▶ It looks like there's some sort of quadratic relationship between horsepower and MPG in the last scatter plot. We want to try and fit a prediction rule of the form

$$H(x) = w_0 + w_1 x + w_2 x^2$$

- ▶ Note that while this is quadratic in horsepower, it is **linear in the parameters!**
- ▶ We can do that, by choosing our two “features” to be  $x_i$  and  $x_i^2$ , respectively.
  - ▶ In other words,  $x_i^{(1)} = x_i$  and  $x_i^{(2)} = x_i^2$ .
  - ▶ More generally, we can create new features out of existing features.

## A quadratic prediction rule

- ▶ Desired prediction rule:  $H(x) = w_0 + w_1x + w_2x^2$ .
- ▶ The resulting design matrix looks like this:

$$X = \begin{bmatrix} 1 & x_1 & x_1^2 \\ 1 & x_2 & x_2^2 \\ \dots & & \\ 1 & x_n & x_n^2 \end{bmatrix}$$

- ▶ To find optimal parameter vector  $\vec{w}^*$ : solve the **normal equations**!

$$X^T X w^* = X^T y$$

## More examples

- ▶ What if we want to use a prediction rule of the form  
$$H(x) = w_0 + w_1 x + w_2 x^2 + w_3 x^3?$$
- ▶ What if we want to use a prediction rule of the form  
$$H(x) = w_1 \frac{1}{x^2} + w_2 \sin x + w_3 e^x?$$

# Feature engineering

- ▶ The process of creating new features out of existing information in our dataset is called **feature engineering**.
  - ▶ In this class, feature engineering will mostly be restricted to creating non-linear functions of existing features (as in the previous example).
  - ▶ In the future you'll learn how to do other things, like encode categorical information.

## Non-linear functions of multiple features

- Recall our example from last lecture of predicting sales from square footage and number of competitors. What if we want a prediction rule of the form

$$\begin{aligned}H(\text{sqft}, \text{comp}) &= w_0 + w_1 \text{sqft} + w_2 \text{sqft}^2 \\&\quad + w_3 \text{comp} + w_4 \text{sqft} \cdot \text{comp} \\&= w_0 + w_1 s + w_2 s^2 + w_3 c + w_4 sc\end{aligned}$$

- Make design matrix:

$$X = \begin{bmatrix} 1 & s_1 & s_1^2 & c_1 & s_1 c_1 \\ 1 & s_2 & s_2^2 & c_2 & s_2 c_2 \\ \dots & \dots & \dots & \dots & \dots \\ 1 & s_n & s_n^2 & c_n & s_n c_n \end{bmatrix}$$

Where  $s_i$  and  $c_i$  are square footage and number of competitors for store  $i$ , respectively.

## Finding the optimal parameter vector, $\vec{w}^*$

- ▶ As long as the form of the prediction rule permits us to write  $\vec{h} = X\vec{w}$  for some  $X$  and  $\vec{w}$ , the mean squared error is

$$R_{\text{sq}}(\vec{w}) = \frac{1}{n} \|\vec{y} - X\vec{w}\|^2$$

- ▶ Regardless of the values of  $X$  and  $\vec{w}$ ,

$$\begin{aligned}\frac{dR_{\text{sq}}}{d\vec{w}} &= 0 \\ \implies -2X^T\vec{y} + 2X^TX\vec{w} &= 0 \\ \implies X^TX\vec{w}^* &= X^T\vec{y}.\end{aligned}$$

- ▶ The **normal equations** still hold true!

# Linear in the parameters

- ▶ We can fit rules like:

$$w_0 + w_1 x + w_2 x^2 \quad w_1 e^{-x^{(1)2}} + w_2 \cos(x^{(2)} + \pi) + w_3 \frac{\log 2x^{(3)}}{x^{(2)}}$$

- ▶ This includes arbitrary polynomials.
- ▶ We can't fit rules like:

$$w_0 + e^{w_1 x} \quad w_0 + \sin(w_1 x^{(1)} + w_2 x^{(2)})$$

- ▶ We can have any number of parameters, as long as our prediction rule is **linear in the parameters**, or linear when we think of it as a function of the parameters.

## Determining function form

- ▶ How do we know what form our prediction rule should take?
- ▶ Sometimes, we know from *theory*, using knowledge about what the variables represent and how they should be related.
- ▶ Other times, we make a guess based on the data.
- ▶ Generally, start with simpler functions first.
  - ▶ Remember, the goal is to find a prediction rule that will generalize well to unseen data.



## Discussion Question

Suppose you collect data on the height, or position, of a freefalling object at various times  $t_i$ . Which form should your prediction rule take to best fit the data?

- a) constant,  $H(t) = w_0$
- b) linear,  $H(t) = w_0 + w_1 t$
- c) quadratic,  $H(t) = w_0 + w_1 t + w_2 t^2$
- d) no way to know without plotting the data

## Example: Amdahl's Law


- ▶ Amdahl's Law relates the runtime of a program on  $p$  processors to the time to do the sequential and nonsequential parts on one processor.

$$H(p) = t_s + \frac{t_{NS}}{p}$$

- ▶ Collect data by timing a program with varying numbers of processors:

Processors	Time (Hours)
1	8
2	4
4	3

**Example: fitting  $H(x) = w_0 + w_1 \cdot \frac{1}{x}$**



$x_i$	$y_i$
1	8
2	4
4	3

## Example: Amdahl's Law

- ▶ The solution is:  $t_S = 1$ ,  $t_{NS} = \frac{48}{7} \approx 6.86$
- ▶ Therefore our prediction rule is:

$$\begin{aligned} H(p) &= t_S + \frac{t_{NS}}{p} \\ &= 1 + \frac{6.86}{p} \end{aligned}$$

# Transformations

# How do we fit prediction rules that aren't linear in the parameters?

- Suppose we want to fit the prediction rule

$$H(x) = w_0 e^{w_1 x}$$

This is **not** linear in terms of  $w_0$  and  $w_1$ , so our results for linear regression don't apply.

- **Possible Solution:** Try to apply a **transformation**.

# Transformations

- ▶ **Question:** Can we re-write  $H(x) = w_0 e^{w_1 x}$  as a prediction rule that **is** linear in the parameters?

## Transformations

- ▶ **Solution:** Create a new prediction rule,  $T(x)$ , with parameters  $b_0$  and  $b_1$ , where  $T(x) = b_0 + b_1 x$ .
  - ▶ This prediction rule is related to  $H(x)$  by the relationship  $T(x) = \log H(x)$ .

- ▶  $\vec{b}$  is related to  $\vec{w}$  by  $b_0 = \log w_0$  and  $b_1 = w_1$ .

- ▶ Our new observation vector,  $\vec{z}$ , is 
$$\begin{bmatrix} \log y_1 \\ \log y_2 \\ \dots \\ \log y_n \end{bmatrix}.$$

- ▶  $T(x) = b_0 + b_1 x$  is linear in its parameters,  $b_0$  and  $b_1$ .
- ▶ Use the solution to the normal equations to find  $\vec{b}^*$ , and the relationship between  $\vec{b}$  and  $\vec{w}$  to find  $\vec{w}^*$ .



# Demo

Let's try this out in a Jupyter notebook. [Follow along here.](#)

## Non-linear prediction rules in general

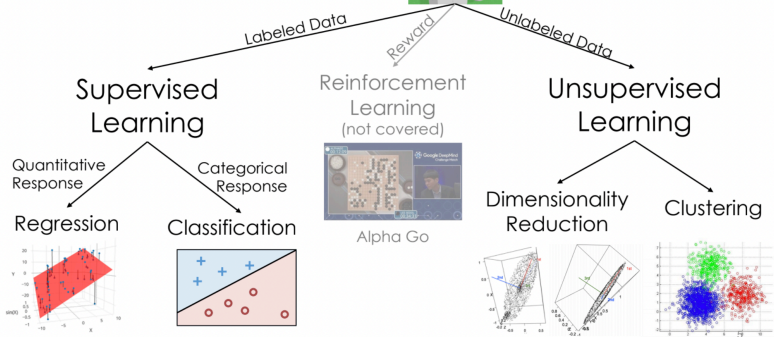
- ▶ Sometimes, it's just not possible to transform a prediction rule to be linear in terms of some parameters.
- ▶ In those cases, you'd have to resort to other methods of finding the optimal parameters.
  - ▶ For example, with  $H(x) = w_0 e^{w_1 x}$ , we could use gradient descent or a similar method to minimize mean squared error,  $R(w_0, w_1) = \frac{1}{n} \sum_{i=1}^n (y_i - w_0 e^{w_1 x_i})^2$ , and find  $w_0^*, w_1^*$  that way.
- ▶ Prediction rules that are linear in the parameters are much easier to work with.

# Taxonomy of machine learning

# What is machine learning?

- ▶ **One definition:** Machine learning is about getting a computer to find patterns in data.
- ▶ Have we been doing machine learning in this class? **Yes.**
  - ▶ Given a dataset containing salaries, predict what my future salary is going to be.
  - ▶ Given a dataset containing years of experience, GPAs, and salaries, predict what my future salary is going to be given my years of experience and GPA.


# Taxonomy of Machine Learning



1

# Clustering

**Question: how might we “cluster” these points into groups?**



## Problem statement: clustering

**Goal:** Given a list of  $n$  data points, stored as vectors in  $\mathbb{R}^d$ ,  $\vec{x}_1, \vec{x}_2, \dots, \vec{x}_n$ , and a positive integer  $k$ , **place the data points into  $k$  groups of nearby points.**

- ▶ These groups are called “clusters”.
- ▶ Think about groups as **colors**.
  - ▶ i.e., the goal of clustering is to assign each point a color, such that points of the same color are close to one another.
- ▶ Note, unlike with regression, there is no “right answer” that we are trying to predict — there is no  $y$ !
  - ▶ Clustering is an **unsupervised** method.





## How do we define a group?

- ▶ One solution: pick  $k$  cluster centers, i.e. **centroids**:

$$\vec{\mu}_1, \vec{\mu}_2, \dots, \vec{\mu}_k \text{ in } \mathbb{R}^d$$

- ▶ These  $k$  centroids define the  $k$  groups.
- ▶ Each data point “belongs” to the group corresponding to the nearest centroid.
- ▶ This reduces our problem from being “find the best group for each data point” to being “find the best locations for the centroids”.



## How do we pick the centroids?

- ▶ Let's come up with an **cost function**,  $C$ , which describes how good a set of centroids is.
  - ▶ Cost functions are a generalization of empirical risk functions.
- ▶ One possible cost function:

$C(\mu_1, \mu_2, \dots, \mu_k)$  = total squared distance of each data point  $\vec{x}_i$  to its closest centroid  $\mu_j$

- ▶ This  $C$  has a special name, **inertia**.
- ▶ Lower values of  $C$  lead to “better” clusterings.
  - ▶ **Goal:** Find the centroids  $\mu_1, \mu_2, \dots, \mu_k$  that minimize  $C$ .

## Discussion Question

Suppose we have  $n$  data points,  $\vec{x}_1, \vec{x}_2, \dots, \vec{x}_n$ , each of which are in  $\mathbb{R}^d$ .

Suppose we want to cluster our dataset into  $k$  clusters. How many ways can we assign points to clusters?

- a)  $d \cdot k$
- b)  $d^k$
- c)  $n^k$
- d)  $k^n$
- e)  $n \cdot k \cdot d$

## How do we minimize inertia?

- ▶ **Problem:** there are exponentially many possible clusterings. It would take too long to try them all.
- ▶ **Another Problem:** we can't use calculus or algebra to minimize  $C$ , since to calculate  $C$  we need to know which points are in which clusters.
- ▶ We need another solution.

# k-Means Clustering, i.e. Lloyd's Algorithm

Here's an algorithm that attempts to minimize inertia:

1. Pick a value of  $k$  and randomly initialize  $k$  centroids.
2. Keep the centroids fixed, and update the groups.
  - ▶ Assign each point to the nearest centroid.
3. Keep the groups fixed, and update the centroids.
  - ▶ Move each centroid to the center of its group.
4. Repeat steps 2 and 3 until the centroids stop changing.

## Example

See the following site for an interactive visualization of k-Means Clustering: <https://tinyurl.com/40akmeans>

**Summary, next time**



## Summary

- ▶ The process of creating new features is called feature engineering.
- ▶ As long as our prediction rule is linear in terms of its parameters  $w_0, w_1, \dots, w_d$ , we can use the solution to the normal equations to find  $\vec{w}^*$ .
  - ▶ Sometimes it's possible to transform a prediction rule into one that is linear in its parameters.
- ▶ Linear regression is a form of supervised machine learning, while clustering is a form of unsupervised learning.
- ▶ Clustering aims to place data points into “groups” of points that are close to one another. k-means clustering is one method for finding clusters.

## Next time

- ▶ How does k-means clustering attempt to minimize inertia?
- ▶ How do we choose good initial centroids?
- ▶ How do we choose the value of  $k$ , the number of clusters?